

-
- Two general types of software: System and Application
- Most important function: process, device, memory and file management
- **System Interface Call:** Used by application programmers to request os services.
- **Process Manager:** Os component that is responsible for creating and destroying processes as well as scheduling them
- **Memory Manager:** Os component that allocated and de-allocates memory and protects against unwanted memory access
- **File Manager:** Os component that creates and deletes files/directories , reads and writes
- **Device Manager:** Os component that uses a device driver to control hardware devices and handles devices
- **Multiprogramming:** When an OS manages several process in memory and shares the CPU
- **Context Switching:** When the OS changes the process currently executing on the CPU
- **Resource Manager:** Allocation and de allocation of resources to components that need it (sharing resources across processes that need them)
- **Process:**
 - **System Process:** Exists to execute OS actions
 - **User Process:** User started process
- **Process State:**

Threads

- **Thread:** A software computational unit
 - Each process has a base thread
 - All threads in one process share code and resources

- **Thread Control Block:** Contains information for each thread
 - Contains an instruction counter and id
 - 1 to 1 app between kernel threads and the app
 - Also can be many to many
- Pool of pre allocated kernel threads are used to service app threads
- **Multithreading:** Process with multiple threads
- **Multiprogramming:** An OS that can manage multiple processes
 - Process can be interrupted for another process (context switching)

Modeling

- **Types of OS Models:**
 - **Physical:** Small scale physical representation
 - **Graphical:** Diagrams
 - **Mathematical:** Algorithms
- **Process Steps:**
 - Memory manager takes program instance at head of queue and allocates memory
 - Processor services program
 - Program terminated, memory is deallocated
- **Way of Measuring System Performance:**
 - Measure the actual system
 - create simulations
 - create analytical models
- **Types of Simulation Models:**
 - **Deterministic:** No randomness in the test
 - **Stochastic:** Has randomness included by having random values inputted

Multiprogramming

- Multiple processes at once
- **Degree of multiprogramming:** Number of processes a system supports

- If there is no memory, the job goes back into the queue
 - If the queue is full the job is lost forever
- **CPU Burst:** Short CPU requests
- **I/O Burst:** Short I/O request

Processor Scheduling

- **Long Term Scheduling:** The OS creates a new process from the job waiting in the input queue
- **Mid Term Scheduling:** The OS decides when and which process to swap in from or to memory
- **Short Term Scheduling:** The OS decides which OS to execute next, also called CPU scheduling
- **CPU Service Requests:**
 - Join the ready queue
 - Execute the CPU burst
 - Join I/O queue
 - Terminate and exit
- **Single Class System:** When a OS treats all processes in the same manner
 - “fair” because there is only one group of processes and they are all treated equally
 - Processes have a priority. Higher priority processes are executed first

- **Synchronization:** coordination of the activities of two or more processes that carry out the following activities
 - Compete for resources in a mutually exclusive manner
 - Cooperate in sequencing or ordering specific events in their individual activities
- **No Synchronization:** When two processes execute without synchronization, which causes problems
- **Mutual Exclusion:** Only one process can use the shared resources at a time. To achieve mutual exclusion:
 - Only one process is executing its critical sections at a time
 - No starvation - Processes wait a finite time interval to enter their critical sections
 - No deadlocks - Processes should not block each other out indefinitely
 - A process will execute its critical section in a finite time interval
- **Critical Sections:** Parts of the code in a process that accesses and uses a shared resource
- **Semaphores:** An abstract data type that functions as a software synchronization tool
 - **Binary Semaphore:** Value can be 0 or 1, used to allow mutual exclusion
 - **Counting Semaphore:** Can be any nonnegative integer value
- **The Bound Buffer Problem:**
 - A producer inserts into buffers, consumer takes out
 - Both go at their own speeds so there is a need to synchronize them
 - Problem: Producers can't add to a full buffer and consumers can't take from an empty buffer so you need to have the buffer be mutually exclusive
 - Solution: Have a counting semaphore to check if the buffer is full, have another to check if the buffer is empty and have a mutex for mutual exclusion
- **The Reader/Writer Problem:**
 - A writer can't access a resource if multiple readers have access so there is a possibility of starvation
- **POSIX Threads:** Each thread has
 - Thread ID
 - A start function
 - An attribute object
 - An argument of the start function
- **Monitors:** A mechanism that implements mutual exclusion

- abstract data types implemented as a class
- **Deadlock:** When the OS is attempting to allocate resources to a set of processes requesting these resources
 - Can cause infinite waiting
 - Conditions for a deadlock:
 - Mutual exclusion: only one process can have a resource
 - Hold and wait: a process holds a resource and requests another
 - Circular wait: a condition that is represented by a closed path of resource allocation and requests
 - No preemption: a process cannot be interrupted and forced to release its resources
 - These are necessary but not sufficient
- **How to handle deadlocks:**
 - Prevention
 - Detection and recovery
 - Avoidance
- **Prevention:**
 - Hold and wait: Two approaches
 - A process must acquire all resources before starting
 - A process must release all resources before getting new ones
 - Circular wait: Circular wait: Assign a total ordering to all resource types
- **Avoidance:** Checks the current resource allocation state to make sure it is a “safe state”
 - The system checks a process “maximum claim” which is the total number of resource instances that a process will request
 - **Bankers Algorithm:** Checks to see if the state is safe
 - A state is defined by the total number of resources, the number of allocated resources and the max resource claim of the process
- **Detection And Recovery:** The OS allocates resources to the process and periodically checks for deadlock
 - **Early detection:** Check after every resource allocation
 - Once detected there are a few options:
 - Terminate all deadlock processes
 - Allocate resources to the deadlocked process
 - Terminate all other processes and run the deadlock process
 - The roll back to before the deadlock

Memory Management:

- **Logical Address:** Reference to some location of a process
- **Process Address:** Set of logical addresses that a process references in its code

- **Symbolic Address:** The address used in a source program
 - ex: Variable names
- **Relative Address:** A compiler converts the symbolic address into a relative address
- **Physical Address:** The final address generated when a program is loaded and ready to execute into memory
- **Partitions:** Blocks of contiguous memory, each one allocated to an active process
- **Fixed Partitions:** Memory is divided into fixed size partitions
 -